

Statistics and graphs

CSW





Susan Stepney

Department of Computer Science
University of York

recap



- Null hypothesis H_0 is assumed true, unless *rejected* by the experimental data
- H_0 = "all swans are white"
 - rejected by a counter-example
 - I see a black swan, and reject H_0
- H_0 = "no effect on readability"
 - rejected by *statistically significant* evidence against
 - I see a certain degree of improvement, and reject H_0
 - reject only at a given *confidence level*
 - how *much* improvement do I need?
 - how *much* confidence can I have?

What do these tell you?



it depends...


- everything varies with the data sample somehow
- is the difference *statistically significant*?



some definitions

- population size N of items $\{x_i\}_{i=1..N}$
- population mean $\mu = \frac{1}{N} \sum_{i=1}^N x_i$
- population standard deviation (RMS) $\sigma = \sqrt{\frac{1}{N} \sum_{i=1}^N (\mu - x_i)^2}$
- sample size n usually $\ll N$ of items $\{x_j\}_{j=1..n}$
- sample mean $m = \frac{1}{n} \sum_{i=1}^n x_i$
- sample standard deviation $s = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (m - x_i)^2}$
 - note the $n - 1$
 - because derived from *estimated* sample mean
 - one fewer "degrees of freedom"

Normal (Gaussian) probability distribution




$$\frac{1}{\sqrt{2\pi}\sigma} \exp\left(\frac{-(x-\mu)^2}{2\sigma^2}\right)$$

standard normal distribution,
 $\mu = 0, \sigma = 1$: $\frac{1}{\sqrt{2\pi}} \exp(-x^2/2)$


- 68.3% lies within ± 1 standard deviation of the mean
- 95.5% lies within ± 2 s.d.
- 99.7% (essentially "all") lies within ± 3 s.d.
 - it is (often incorrectly!) assumed that the underlying distribution is normal
 - there are certain tricks if it isn't

normal distributions ; $\mu = \pm 10$; $n = 5$


$\sigma = 5$



$\sigma = 10$




$\sigma = 50$




[inspired by Wineberg & Christensen, tutorial, CEC 2004]

normal distributions ; $\mu = \pm 10$; $n = 100$


$\sigma = 5$



$\sigma = 10$



$\sigma = 50$




estimating the true mean from a sample

- started with a normal population, mean μ and s.d. σ
- generated *samples* drawn from this, $n = 5$, $n = 100$
- calculated the sample mean m and s.d. s
 - the sample mean is different from the true mean
 - can we use it to *estimate* the true mean?
 - because, in real cases, we *don't know* the true mean!
- intuitively:
 - the smaller the standard deviation, the closer to the true mean
 - the bigger the sample, the closer to the true mean

t distribution


- we have a sample with mean m , from a population with mean μ
 - *statistics*: take samples from a standard normal distribution, then the sample means m_i have a " t distribution"
 - family of distributions, depending (weakly) on n
 - area under a region of the curve is the probability that m lies in that region
- also called the "Student t distribution", after the pen-name of statistician W. S. Gossett

$$t = \frac{m - \mu}{s / \sqrt{n}}$$



confidence levels

- “95% confidence” : only $\alpha = 5\%$ of the area under the curve lies in the tail(s)
 - only 5% chance that the mean lies outside the bounds
 - “99% confidence” : $\alpha = 1\%$ lies in the tail(s)
- books of tables of the value of t_0 for different confidence levels α and different sample sizes n
 - one tailed test: $t_{\alpha}(n - 1)$
 - two tailed test: $t_{\alpha/2}(n - 1)$
- or MS-Excel function
 - one tailed test: **TINV(2* α , $n - 1$)**
 - two tailed test: **TINV(α , $n - 1$)**



confidence levels (two tailed) $\mu = 10, \sigma = 10$

- $n = 5$

measured mean $m = 17.0$, standard deviation $s = 5.4$

$$t_0 = \frac{m - \mu_0}{s / \sqrt{n}}$$

95% confidence (1 in 20 chance of being wrong):

$TINV(0.05, 4) = 2.78$; $\mu = m \pm 2.78 * s / \text{root } n$

$\mu = 17.0 \pm 6.7$; $10.3 \leq \mu \leq 23.7$

$$\mu_0 = m \pm t_0 \frac{s}{\sqrt{n}}$$

99% confidence: $\mu = 17.0 \pm 11.2$; $5.8 \leq \mu \leq 28.2$

99.9% confidence: $\mu = 17.0 \pm 20.9$; $-3.9 \leq \mu \leq 37.9$

- $n = 100$

measured mean $m = 9.0$, standard deviation $s = 10.3$

95% confidence: $\mu = 9.0 \pm 2.1$; $6.9 \leq \mu \leq 11.1$

99% confidence: $\mu = 9.0 \pm 2.7$; $6.3 \leq \mu \leq 11.7$

99.9% confidence: $\mu = 9.0 \pm 3.5$; $5.5 \leq \mu \leq 12.5$


are two means different?

- use the (Student) t test, to calculate the probability p that two sample means m_1 and m_2 are the same
- calculate the t statistic
$$t_0 = \frac{m_1 - m_2}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}}$$
- calculate the probability that the means are same, which is the area under the t distribution between $[-t_0, t_0]$
 - calculate in MS-Excel: $p = \text{TDIST}(t_0, n_1 + n_2 - 2, 2)$


confidence levels (two tailed) $\mu = \pm 10, \sigma = 50$

- $n = 5$
 - measured : $m_1 = -11.5, m_2 = -28.9, s_1 = 28.7, s_2 = 55.5$
 $t_0 = 0.624$; $p = \text{TDIST}(0.624, 8, 2) = 0.55$
 - probability the two means are *same* = **55%**
 - would *you* bet money on it?
- $n = 100$
 - measured : $m_1 = 16.3, m_2 = -12.6, s_1 = 44.9, s_2 = 51.5$
 $t_0 = 4.23$; $p = \text{TDIST}(4.23, 198, 2) = 3.5 \times 10^{-5}$
 - probability the two means are *same* = **0.0035%**
 - I'd put money on them being different !

reminder : $\mu = \pm 10$; $\sigma = 50$; $n = 5, 100$



$p(\text{means different}) = 45\%$



$p(\text{means different}) = 1$

"significant" does not mean "important"

- the bigger the samples, the better the statistical significance - that must be good?
 - can nearly *always* get a statistically significant result just by having a big enough sample size !
 - with *enough* samples, can distinguish between two means ...
 - but it might not be an *important* difference
 - ... but the two means might be very *very* close ...
 - the old algorithm has a mean success rate of 52.38%
 - whereas *my* algorithm's success rate is 52.41%
 - improvement significant at the 99.9% confidence level
 - so why aren't you impressed by my result?
 - ... because it's a very small *effect*
 - happens easily when experimental runs are "cheap"

effect size : Cohen's d

- measure of *importance*

- here assume $n_1 = n_2 = n$

- 0.2** : a small effect

- the data's *dispersion* is bigger than the difference in the means

- 0.5 : a medium effect

- 0.8** : a big effect

- may be worth getting excited about

- example: $\sigma = 50, n = 100, t_0 = 4.23$

- probability that means are the same = **0.0035%**

- but has $d = 0.6$: so don't get *too* excited by it

- example: $\sigma = 10, n = 100, t_0 = 14.2$

- probability that means are the same = **0.0%**

- and has $d = 2.0$

- so it's a big effect, too

$$d = t \sqrt{\frac{2}{n-1}}$$

[J. Cohen. *Statistical power analysis for the behavioral sciences*, 2nd edn. 1988.
<http://web.uccs.edu/lbecker/Psy590/es.htm>]

non-parametric tests

- all this so far assumes a normal probability distribution
 - some distributions are very affected by far “outliers”
 - long “tails” can dominate the mean, and make the standard deviation meaningless
 - there are tricks you can play if your distribution isn't normal
 - Central Limit Theorem, ...
- can use *non-parametric tests*
 - “no parameters” - no assumption about the *shape* of the probability distribution function
 - uses *rank ordering* instead of values
 - intuition: a lowest ranked outlier sample is just “the worst”; is doesn't matter how *much* worse than the rest it is

median : the “non-parametric mean”

- mean : item with average *value*
 - $\text{mean}\{-20, 1, 2, 3, 4\} = -2$
 - $\text{mean}\{-20, 1, 2, 3\} = -3.5$
- median : item with average *rank*
 - rank the items in order, and pick the middle one
 - $\text{median}\{-20, 1, 2, 3, 4\} = 2$
 - if there is an even number of data items, average the two values
 - $\text{median}\{-20, 1, 2, 3\} = 1.5$
- median = 50th percentile
 - quartiles (25th and 75th percentiles) are a non-parametric measure of *spread*


are two medians different?

- there are non-parametric versions of the t test to check if two ranked samples are significantly different
- also analogues of confidence tests, etc
- if you need to use these - check with an expert!


is your picture worth a kilo-word?

CS	37
Electronics	44
Maths	10
English	9


best !
for *this* data
(get more data?)




Chartjunk




small distinctions obscured



More Chartjunk



Cluttered


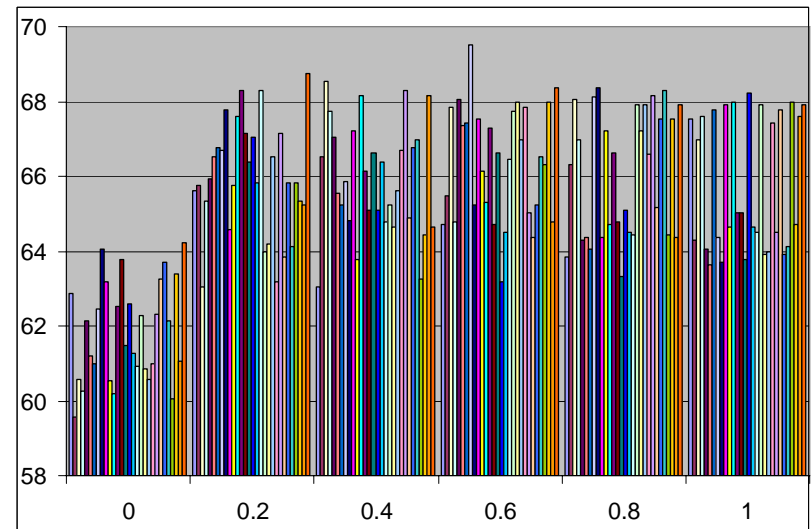


better

[Tufte. *The Visual Display of Quantitative Information*. Graphics Press, 1983]

plot your data to expose its structure

0.00	0.20	0.40	0.60	0.80	1.00
62.86	65.61	63.06	64.70	63.85	67.54
59.58	65.74	66.51	65.48	66.32	64.31
60.59	63.06	68.55	67.86	68.05	66.96
60.26	65.35	67.73	64.77	66.96	67.60
62.13	65.93	67.03	68.05	64.31	64.05
61.20	66.51	65.55	67.35	64.38	63.65
61.00	66.77	65.22	67.41	64.05	67.79
62.46	66.71	65.87	69.50	68.11	64.38
64.05	67.79	64.83	65.22	68.36	63.72
63.19	64.57	67.22	67.54	64.38	67.92
60.53	65.74	63.79	66.13	67.22	64.64
60.19	67.60	68.17	65.29	64.70	67.98
62.53	68.30	66.13	67.28	66.64	65.03
63.79	67.16	65.09	64.70	64.77	65.03
61.47	66.39	66.64	66.64	63.32	63.79
62.60	67.03	65.09	63.19	65.09	68.24
61.26	65.81	66.39	64.51	64.51	64.64
60.93	68.30	64.77	66.45	64.44	64.51
62.27	63.98	65.22	67.73	67.92	67.92
60.86	64.18	64.64	67.98	67.22	63.92
60.59	66.51	65.61	66.96	67.92	63.98
61.00	63.19	66.71	67.86	66.58	67.41
62.33	67.16	68.30	65.03	68.17	64.51
63.26	63.85	64.90	64.38	65.16	67.79
63.72	65.81	66.77	65.22	67.54	63.92
62.13	64.11	66.96	66.51	68.30	64.11
60.05	65.81	63.26	66.32	64.44	67.98
63.39	65.35	64.44	67.98	67.54	64.70
61.06	65.22	68.17	64.77	64.38	67.60
64.24	68.74	64.64	68.36	67.92	67.92



[Clark, Jacob, Stepney. *Secret Agents Leave Big Footprints. GECCO 2003*]

plot all your data, to see outliers



- you have some “anomalous” data



- don't just discard it as an “outlier” - *understand* it!
 - is it just a statistical fluctuation?
 - a once-in-a-blue-moon “six sigma” outlier?
 - is it an error in the experimental design or implementation?
 - fix the problem, and rerun *all* the experiments
 - is it in interesting unexpected effect?
 - investigate it further!
 - it *might* be the basis of a new discovery

3-sigma error bars



- can use a *scatterplot* to show *all* the data
- can also *summarise* the data with a *statistic*
 - show spread as well as average, with *error bars*
 - normal "3-sigma" error bars: encompass ~ 99.7% of the data




"box-and-whisker" error bars

- median, quartiles

- John W. Tukey. *Exploratory data analysis*. Addison Wesley, 1977



"deluxe" box-and-whisker bars




use a "cutoff" to highlight outliers

plot *mean*, to highlight skew

"notches" at
 $\text{median} \pm 1.58(\text{IQR} / \sqrt{n})$

- if notches on separate bars do not overlap, ~ 95% confidence the medians are different
- small samples may have "folded back" notches outside the IQR:



summary : statistics

- don't present single data points, or even just the mean
 - show the data *spread*
- do use t distribution to calculate *confidence levels*
 - try to use *at least* the 99% level, preferably 99.9%
- do calculate *importance* as well as *significance*
- do plot your data
 - *with error bars ; without chartjunk*
- *lots* more to experimental design, statistics and graphs
 - other measures of fit
 - "factorial" designs for controlling multiple variables
 - "small multiples", "stem-and-leaf" plots, ...